

Token-level Recall (%)

Head Retention (%)

Retention

Memory (GB)

Elbow Point:
Prune Ratio: 83.7%
Recall: 21.33%
Speedup: 1.31x
Retention: 1.007

- Token-level Recall
- Dense Model
- Elbow Point
- Light Pruning
- Heavy Pruning
- Retention
- Min Speedup (0.43x)
- Max Speedup (1.67x)

